A/B Testing: The Most Powerful Way to Turn Clicks into Customers
By Dan Siroker, Pete Koomen and Cara Harshman

CHAPTER

12

Beyond the Page: Non-Website A/B Testing

How to Test Email and Pricing

Up to this point, we've primarily looked at examples where teams test a page's various design elements: forms, images, headlines, and layout. This chapter explores A/B testing beyond the page, in two different domains: email and pricing.

The *What* and the *When*: Prezi

Email is one of the most important things you can A/B test. Just as you can use A/B testing to make the elements of your website work better, you can A/B test emails to increase quantifiable success metrics like *open rate* and *click-through rate*. And just as you can roll out a website change gradually and gauge response before showing it to all users, it's easy and incredibly advantageous to roll out an email in the same way. First, select a portion of your total mailing list and send them a number of variations. Gauge your success metrics, and *then* send the email that works best to the remainder of your full list.

It was Halloween of 2011, and David Malpass—marketing analyst at cloud-based presentation tool Prezi—was planning to send out a huge email newsletter blast. He sent his boss, the director of marketing, three candidate subject lines:

1. "New Tricks & Treats"
2. "New features: Templates, Google Image Search"
3. "Create a beautiful Prezi in minutes"

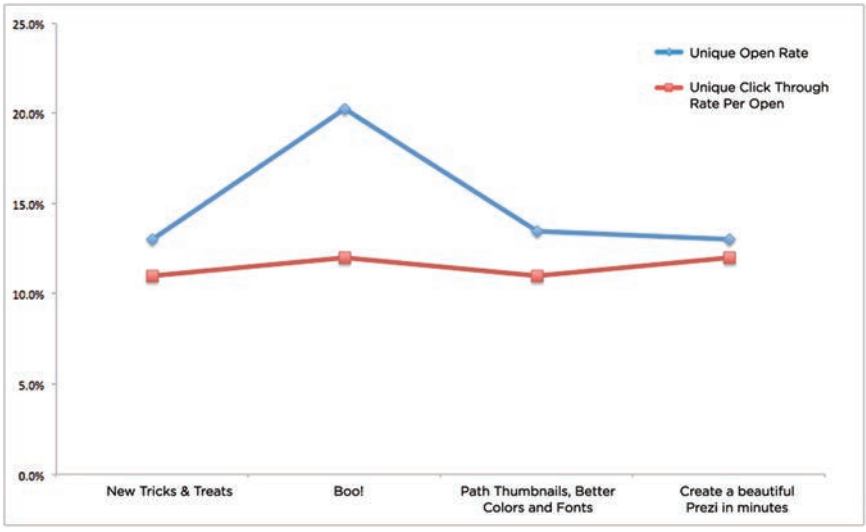

FIGURE 12.1 The open rates for the four Prezi email subject lines.
Source: Prezi

Malpass recalls his boss's response: "Why don't you try one that just says '*Boo!*'?" Malpass thought it was a terrible idea.

Malpass was planning to A/B test the subject line anyway, so he reluctantly threw "*Boo!*" into the mix and sent out emails with the four different subject lines to about a million users. He remembers the results with a smile: "*Boo!*" trounced the other three subject lines with a whopping 20 percent more opens (Figure 12.1).

Email also offers some intriguing opportunities for tests that don't have anything to do with either the content *or* the formatting. Namely, because email (unlike a website) *actively* reaches users (instead of passively awaiting their arrival), the *timing* of the message can itself be tested.

As it turns out, the *hour of the day* and the *day of the week* can both matter in a big way. Educational licenses make up a large portion of Prezi's user base, and Malpass was in for a surprise when he tested for the day of the week with the optimal open rate. He got

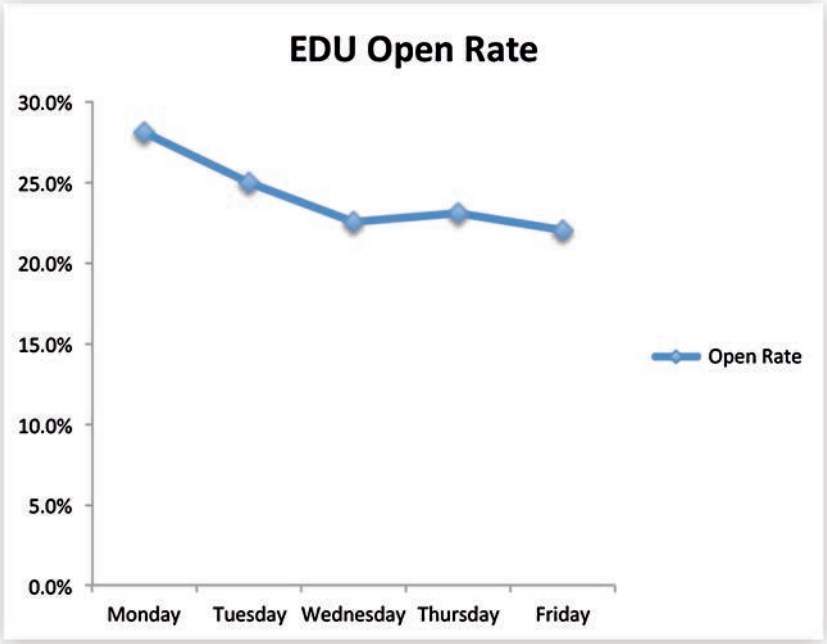

FIGURE 12.2 Open rate graph for educational licenses.

not one answer, but *two*: for educators, Monday had the highest open rate, and the rate decreased through the week (Figure 12.2). For non-educators, open rate peaked *late* in the week, on Thursday (Figure 12.3).

Price Testing

There's another pivotal element to every web business that's asking to be tested, beyond the layout and media and copy and even newsletter timing, and arguably even more important: the price. *Price testing* is a very valuable way to understand how demand for your product changes with price increases and decreases.

"Price testing is some of the most valuable testing and some of the most challenging technically," explains Scott Zakrajsek,

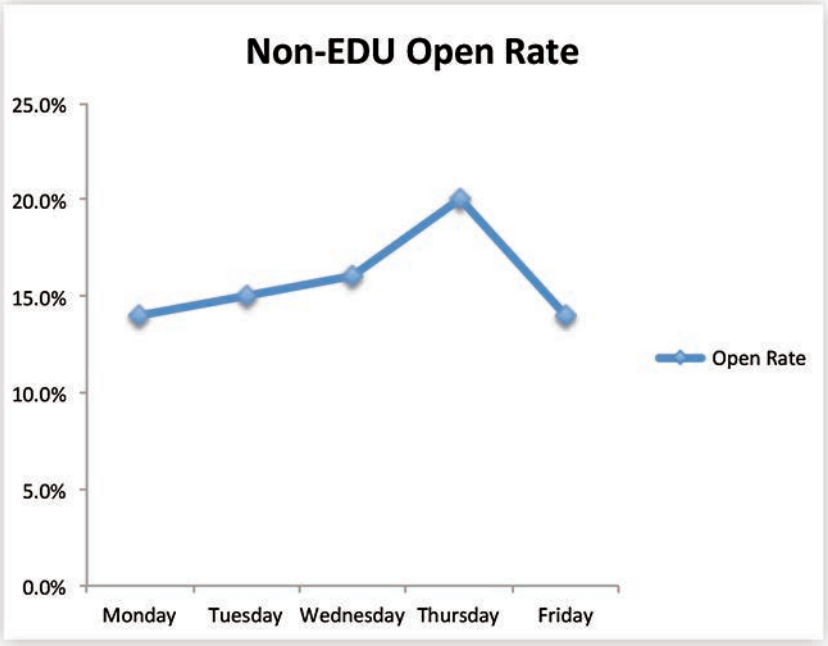

FIGURE 12.3 Open rate graph for non-educational licenses.

who has done testing at e-commerce sites like Adidas, Victoria's Secret, and Staples. "If you're a smaller site with ten to twenty key products, then price testing is the number-one thing you can do to maximize your conversion, your revenue."

The vast majority of prices today aren't really well thought through. It's easy to determine *cost-plus pricing*, where you take your own costs of production and add some kind of profit margin, as well as *competitive* or *market-based* pricing, but much trickier to come up with *value-based pricing*—getting the price in line with the value that you deliver.

With A/B testing, conventional wisdom and generalizations surrounding pricing become obsolete. There are many services out

there where if you increase the price, just as many people will buy it. In some rare cases, and especially if your customers are buying something on behalf of their company, they are *more likely* to buy if the price is higher. Your customers might ask: if other businesses are paying this much for it, it must be good, right? Every product's and service's demand curve will look different, and *testing* prices is one of the best ways to get data about your own.

There's just as much theory and anecdote surrounding a price's trailing digits: the cents. Some people argue that round-dollar figures connote a premium product (and have design possibilities that include eliminating the decimal portion altogether); others argue that there is an important lure to prices that end in 99¢. Some authorities will argue the differences between 95¢, 96¢, 97¢, and 98¢; Walmart is known for its seemingly random cent-pricing ($6.37 toothpaste, $8.02 stapler), which seems to connote that prices have been cut to the bone. Where does that leave you and *your* business?

What's interesting is that these are all psychological tactics that are very difficult to understand fully or to generalize from another business's results to your own. By A/B testing them you can very quickly measure and see for yourself *exactly* how those factors are in play with your own business. You can plot your own price elasticity curve, and in some cases you'll see that if you increase the price, consumption stays the same. The only way to know for certain is to test it.

Testing the Perceived Price: The Last-Minute Discount

There are two hurdles to successful price testing. The first is technical: most prices are stored in a database or pricing table, and

so any test that modifies the price of an item is going to need much *deeper technical integration* than a test simply involving only front-end design elements. The second hurdle is about managing customer relationships: to accidentally show users a lower price on the page and then charge them a higher price breaks not only customers' trust but the law. And users who are charged a higher price even if the site is transparent about it may feel betrayed or alienated when they learn that most other customers paid less.

There is a quick-and-easy method to foray into price testing that avoids *all* of these problems, something we call the *last-minute discount*. The way it works is quite simple. You show the customer a higher price (say $13.99) than the "actual" price of that item, which most users see (say $9.99) and which lives in the back-end price book. As the user in the $13.99 variation group moves down the checkout funnel, they discover *just before* confirming their purchase that a "last-minute discount" has been applied and that their total is in fact only $9.99! You've managed to get reasonably accurate information about what kind of conversion rate to expect from the higher price without dealing with back-end integration *or* worrying about making customers upset.

This type of cosmetic price testing is also perfect for testing the cents of a price. A product shown for $19.63, $19.95, and $19.99 in different tests could check out at $19.50 or $19.00 even; the users may not even notice the change. (Just be sure not to charge customers *more* than they're expecting!) Reading the conventional wisdom about cent pricing yields various—often mutually exclusive—theories and superstitions. The truth is that every business is different; you won't know until you test.

Anchoring in Action: Judy's Book Club

Another way to get data on your pricing without literally offering a product for different prices to different users is to test the way that product's price is *anchored*, or contextualized, for the customer. Consumers typically don't buy the most expensive options available, and many times they don't buy the least expensive, either. (Restaurants, for instance, will often mark up the *second*-least-expensive menu option.) Perhaps displaying only paid plans on your site will get users in at the lowest rung; perhaps offering a "Free" plan with virtually no features will help persuade users that adequate services can't be had for free. Try both.

Judy's Book is a "social search" tool and reviews website that caters to a family audience. Paid business listings are a primary way Judy's Book makes money: businesses get better search positioning and more robust profiles with photos if they pay a monthly fee. Judy's Book wanted to increase the number of businesses that sign up for paid listings, and General Manager Ali Alami hypothesized that positioning a column showing the few features included with a free listing alongside the many features included in the paid listings would increase signups for paid listings (Figure 12.4).

Showing a free listing column increased clicks on the signup button for the basic paid listing by 198.6 percent.

There are a number of additional anchoring techniques worth considering, and, of course, testing. For instance, try adding a tier *above* your most expensive tier in order to make the one beneath it seem less expensive. And try having a small price delta between the plan you *want* people to buy and the one just below it, with the goal being to get customers to think they're getting a bargain ("For just a *little* bit more, I get *all* this stuff!").

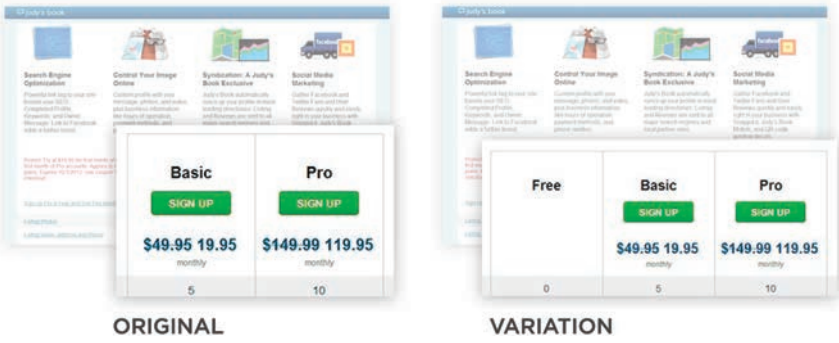

FIGURE 12.4 Two-column Judy's Book listing page versus three-column Judy's Book listing page.

Testing the Billing

You have options when it comes to how the price is broken down to the user as well. You might test representing the price as an annual subscription price or as a monthly cost. You might, for instance, have a monthly price prominently displayed but explain to the user that the product is in fact billed annually. (Just be sure it's clear how much customers will be billed so they don't get a surprise charge.) Testing will reveal how these differences affect conversion, average order value, churn, and so forth.

In testing their own pricing display, for instance, Prezi saw a 12.5 percent increase in signups from emphasizing the monthly breakdown of its annual cost (Figure 12.5).

It's worth keeping *localization effects* in mind here: there are markets like Brazil, for instance, where appliances and hardware are more typically bought with monthly payments than they are in countries like the United States. (Laws can also differ from country to country in ways that affect pricing.) Understanding your market will help you get your bearings, but as always, test and see what works best.

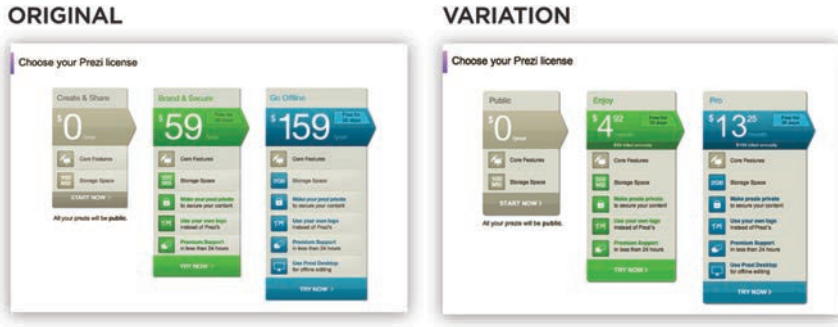

FIGURE 12.5 Prezi pricing displayed annually versus monthly.

Testing the Actual Price

Changing the *actual* price of a good or a service is the most complex of these approaches, in that it requires you to actually change your price book or to add additional SKUs. Full-fledged price testing has a world of nuances and best practices all its own. For instance, perhaps *serial testing* across *all* users at different *times* will avoid potential PR fallout rather than having different prices in play simultaneously. This method has its drawbacks, though: for instance, it becomes difficult to control for the general outside dynamics. One reason why you typically do A/B testing over *different* users in the *same* time period is to control for time-based effects, like day-of-week, time-of-day, news cycle, and the like. Typically, making meaningful sense of serialized price changes requires historical data that can put small-scale fluctuations into context.

A quick final reminder: it's absolutely critical when price testing to make sure you are *defining your success metrics* correctly. A higher price is very likely going to reduce conversions, but increase average order value: the key question is *by how much?* Make sure that you're using an appropriate metric, for instance, *revenue per visitor (RPV)*, which takes into account both

conversion *and* average order value, in order to get accurate reporting from your testing tool about which variation is the winner. Consider, also, that even RPV may not tell the whole story. Higher pricing may lead to customer churn in the medium-term, so think carefully about what your results truly mean before making a major change.

TL;DR

- Not only the **subject lines** but the **timing** can be critical in influencing open rates and click-through rates for email. Roll out large email campaigns gradually, test variations, and then send the winner on to the rest of your list.
- **Price testing** can be one of the most important types of testing to conduct for your business, and there are a number of ways to do it.
- The **last-minute discount** is a great technique for testing price without needing to deal with back-end integration or worrying about upsetting users.
- **Anchoring** the price of a product into context with other prices can greatly affect how users react.
- **Presentation** can be everything when it comes to pricing. For a quick-and-easy price test, try various breakdowns (e.g., monthly, yearly, or weekly) and see what works best.
- **Serial testing** is one way to test prices without needing to show different users different prices at the same time; however, this advantage is offset by difficulties in ensuring the accuracy of its results.